



Chapter 1

Introduction

COMP3314

Machine Learning

Outline

- Motivation
- Types of ML
 - Supervised Learning
 - Classification
 - Regression
 - Reinforcement Learning
 - Chess
 - Unsupervised Learning
 - Clustering
 - Dimensionality Reduction
- Terminology and Notation
- Roadmap
 - Preprocessing
 - Learning
 - Evaluation and Prediction
- Python
 - Installation
- Linear Algebra Review
- References

Motivation

- Nowadays large amount of structured and unstructured data is available
- ML algorithms can turn this data into knowledge
 - Powerful open source libraries available to do this
- In this course you will understand how these algorithms work
- You will also learn how to utilize them to make predictions

Motivation

- ML algorithms are self-learning
 - Automatically derive knowledge from data to make predictions
 - No need for humans to manually derive rules
 - ML offers a more efficient alternative for capturing the knowledge in data to gradually improve the performance of predictive models
- ML becomes increasingly relevant in CS research
 - More importantly
 - Plays an ever greater role in our everyday lives

How do you use machine learning everyday?

Examples of Machine Learning

- Basket analysis
- Credit scoring
- Medical diagnosis
- Biometrics
- Object recognition

Machine Learning Definition

- Subfield of Artificial Intelligence (AI)
- Arthur Samuel (1959)
 - Field of study that gives computers the ability to learn without being explicitly programmed
- Tom Mitchell (1998)
 - A computer program is said to learn from experience E with respect to some task T and some performance measure P , if its performance on T , as measured by P , improves with experience E

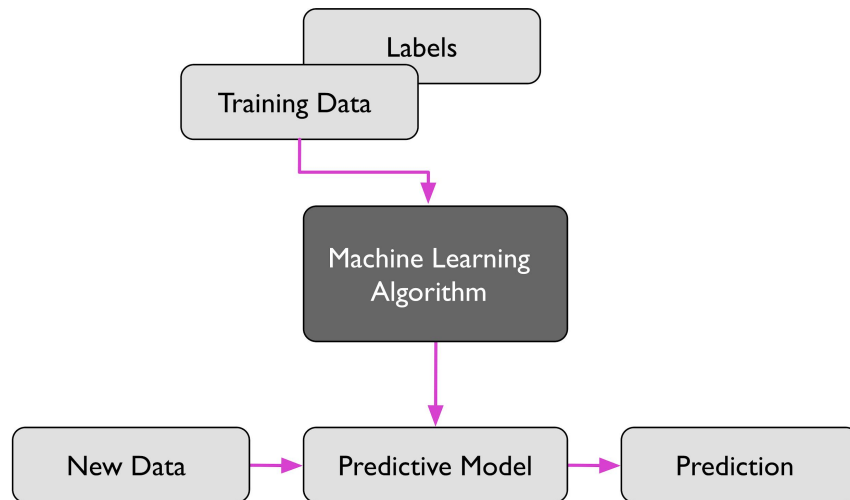
Types of machine learning

- In the following we will consider three types of machine learning

Supervised Learning	<ul style="list-style-type: none">➤ Labeled data➤ Direct feedback➤ Predict outcome/future
Unsupervised Learning	<ul style="list-style-type: none">➤ No labels/targets➤ No feedback➤ Find hidden structure in data
Reinforcement Learning	<ul style="list-style-type: none">➤ Decision process➤ Reward system➤ Learn series of actions

Supervised Learning

- Learn from labeled training data
 - Make predictions about unseen / future data
- Supervised refers to a set of samples where the desired output signals (labels) are already known



Supervised Learning:

Classification vs. Regression


- Two subcategories of supervised learning
 - Classification
 - A supervised learning task with discrete class labels
 - E.g., spam email classifier
 - Regression
 - Outcome is a continuous value
 - E.g., student exam score prediction

Supervised Learning - Classification

- Goal: Predict class labels of new instances, based on past observations
- Class labels are discrete, unordered values
- Two subcategories of classifiers:
 - Binary classification
 - Only two possible class labels can be assigned
 - E.g., spam vs. non-spam emails
 - Multiclass classification
 - Any fixed number >2 of class labels can be assigned
 - E.g., handwritten digit recognition

Classification - Example

- Given 30 training samples
 - 15 labeled as negative class
 - 15 labeled as positive class
- Let each sample have 2 dimensions
- Classifier will learn the decision boundary
 - Represented as a dashed line
 - Able to separate the two classes




Regression

- Prediction of continuous outcome
 - The term regression was devised by Francis Galton in his article [Regression towards Mediocrity](#) in 1886
- Example:
 - Predicting the exam scores given time spent studying

Regression - Example

- Given
 - Predictor variable x
 - Response variable y
- I.e., 1D data set
- Fit a line to it minimizing the distance between sample points and the fitted line
 - Average squared distance is most commonly used
- Use the line to predict outcome of new data

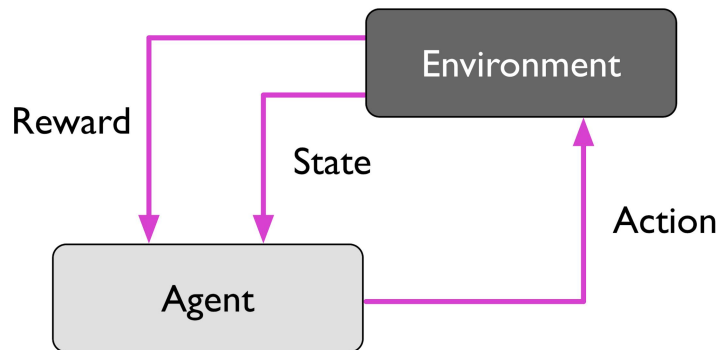


Quiz

- Consider the following supervised ML tasks. Label each task with *Classification Task* or *Regression Task*
 - a. You are working for an investment bank and your task is to predict investors sentiment for certain stocks by analyzing popular online investments forums
 - b. You are working for a property agency and your task is to predict the housing price for a property based on past data that the agency has available in their database
 - c. Your task is to analyze a video stream of the western harbour tunnel and count how many Tesla pass by every day

Reinforcement Learning

- The system (aka agent) improves its performance based on interactions with an environment
- Trial-and-Error approach
 - Learning by doing
- The agent receives feedback (reward) from the environment
 - This reward is not the correct ground truth
 - It is a sample experience
 - Extensive interaction with the environment allows agent to learn a series of actions that maximizes this reward



Reinforcement Learning - Example: Chess


- Agent decides upon a series of moves depending on state of board
 - Environment is the board
 - Reward can be defined as win or lose at the end of the game
- Outcome of each move results in different state of the environment
 - Removing an opponent's chess piece from the board or threatening the queen is associated with a positive event
 - Losing a chess piece to the opponent is associated with a negative event
- Note: Not every turn results in the removal of a chess piece
 - Reinforcement learning is concerned with learning the series of steps by maximizing a reward based on immediate and delayed feedback

Unsupervised Learning

- Unlabeled data / data of unknown structure
- Explores the structure of data
 - Extract meaningful information without guidance of known outcome variable / reward function
- Examples
 - Clustering
 - Dimensionality reduction

Clustering

- Exploratory data analysis technique
- Organizes information into meaningful subgroups (clusters) without having any knowledge of group memberships
- Each cluster defines a group of objects that share a certain degree of similarity but are more dissimilar to objects in other clusters



Clustering - Example

President Trump suggests taking hydroxychloroquine to protect against Covid-19 despite lack of proof

CNN · 1 hour ago



- Trump Again Promotes Hydroxychloroquine For Coronavirus Treatment

NPR · 1 hour ago

- Trump and officials warn the worst is yet to come

CBS Evening News · 4 hours ago

- Andy Puzder In coronavirus crisis, Trump displays leadership Americans expect and want

Fox News · Yesterday · Opinion

Rectangular Snip

- Charlie Kirk: Time for President Trump's New Wollman Rink Moment | Opinion

Newsweek · Yesterday · Opinion

[View Full Coverage](#)

Full Coverage

Sort

Boris Johnson admitted to hospital with coronavirus

Top coverage

SLATE

British Prime Minister Boris Johnson Admitted to Hospital for Coronavirus

2 hours ago



Fox News

Boris Johnson hospitalized after experiencing coronavirus symptoms, PM's office says

7 hours ago



CNN

Boris Johnson admitted to hospital for tests. He has previously tested positive for coronavirus

2 hours ago



Callitman

PM's Covid-19 timeline: from 'mild symptoms' to hospital admission

4 hours ago



Clustering - More Examples


- Human genetic clustering
- Sequence clustering
- Social network analysis
- Market research
- Grouping of shopping items

Dimensionality Reduction

- Often we are working with data of high dimensionality
 - I.e., each observation comes with a high number of measurements
- High dimensional data can present a challenge
 - Computational performance
 - Predictive performance
 - Visualization
- Dimensionality reduction is a commonly used approach in feature preprocessing
 - Compress data onto a smaller dimensional subspace
 - Retaining most relevant information

Dimensionality Reduction - Example

- High-dimensional feature set can be projected onto 1D, 2D or 3D feature spaces
 - 3D to 2D example



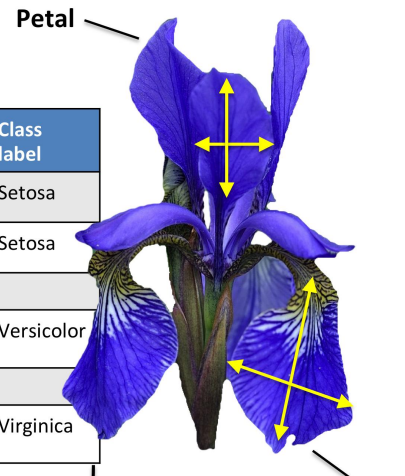
Types of machine learning

- In the following we will consider three types of machine learning

Supervised Learning	<ul style="list-style-type: none">➤ Labeled data➤ Direct feedback➤ Predict outcome/future
Unsupervised Learning	<ul style="list-style-type: none">➤ No labels/targets➤ No feedback➤ Find hidden structure in data
Reinforcement Learning	<ul style="list-style-type: none">➤ Decision process➤ Reward system➤ Learn series of actions

Terminology and Notations

- Iris flower data set contains measurements of 150 Iris flowers from three different species
 - Setosa, Versicolor, and Virginica
- Introduced in Fisher's 1936 paper
The use of multiple measurements in taxonomic problems
- Row
 - A single flower sample
- Column
 - Flower features (measurements in centimeters)



The diagram shows a blue Iris flower with yellow arrows indicating measurements. One arrow points to the length of a petal, another to the width of a sepal, and a third to the length of a sepal.

Samples (instances, observations)					
	Sepal length	Sepal width	Petal length	Petal width	Class label
1	5.1	3.5	1.4	0.2	Setosa
2	4.9	3.0	1.4	0.2	Setosa
...					
50	6.4	3.5	4.5	1.2	Versicolor
...					
150	5.9	3.0	5.0	1.8	Virginica

Features
(attributes, measurements, dimensions)

Class labels
(targets)

Terminology and Notations

- We will use a matrix and vector notation to refer to our data
- Each sample is a separate row in a feature matrix \mathbf{X} , where each feature is stored as a separate column
- Iris dataset example
 - 150 samples and four features are written as a 150 x 4 matrix \mathbf{X}

$$\begin{bmatrix} x_1^{(1)} & x_2^{(1)} & x_3^{(1)} & x_4^{(1)} \\ x_1^{(2)} & x_2^{(2)} & x_3^{(2)} & x_4^{(2)} \\ \vdots & \vdots & \vdots & \vdots \\ x_1^{(150)} & x_2^{(150)} & x_3^{(150)} & x_4^{(150)} \end{bmatrix}$$

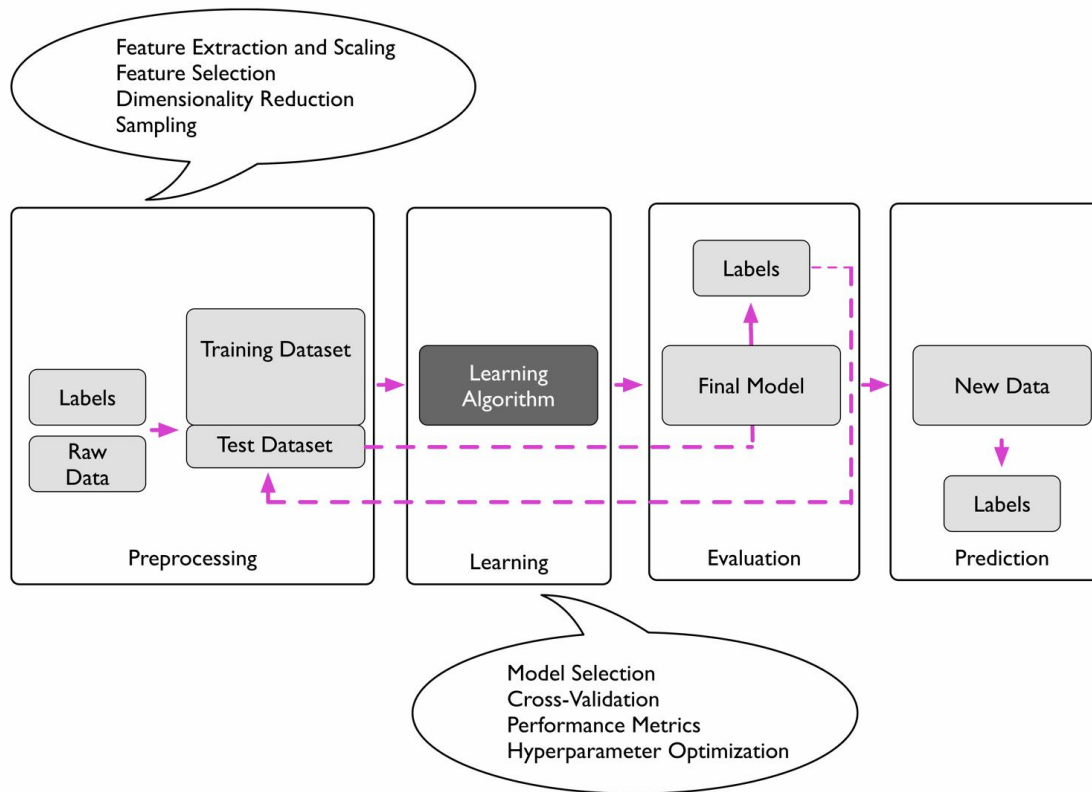
Terminology and Notations

- We will use the superscript i to refer to the i th training sample, and the subscript j to refer to the j th dimension of the dataset
 - For example $x_j^{(i)} = x_1^{(150)}$ refers to the first dimension of the flower sample 150
- We use lowercase, bold-face letters to refer to vectors and uppercase, bold-face letters to refer to matrices
- Note that each row in the iris dataset \mathbf{X} can be written as a four-dimensional row vector and each feature dimension is a 150-dimensional column vector

$$\mathbf{x}^{(i)} = \begin{bmatrix} x_1^{(i)} & x_2^{(i)} & x_3^{(i)} & x_4^{(i)} \end{bmatrix} \quad \mathbf{x}_j = \begin{bmatrix} x_j^{(1)} \\ x_j^{(2)} \\ \vdots \\ x_j^{(150)} \end{bmatrix}$$

Roadmap

- Typical workflow for using ML in predictive modeling



Preprocessing

- Preprocessing of the data is a crucial steps in any ML application
- Feature selection, extraction and scaling
 - Select and extract useful features from raw data
 - Many algorithms also require that the selected features are on the same scale
- Dimensionality reduction
 - May improve
 - Computational performance
 - Predictive performance
- Sampling
 - Randomly divide the dataset into a separate training and test set to determine whether our algorithm not only performs well on the training set but also generalizes well to new data
 - Keep the test set until the very end to evaluate the final model

Learning


- Model selection
 - Compare algorithms and select the best performing model
- Cross-validation
 - How do we know which model performs well on the final test dataset if we don't use this test set for model selection?
 - Cross-validation splits the training dataset further into training and validation subsets
- Performance metric
 - Decide upon a metric to measure performance
- Hyperparameter optimization
 - Fine-tune parameters of the model based on performance on validation set

Evaluation and Prediction

- After model selection and training we use the test dataset to estimate how well it performs on unseen data
 - Estimate the generalization error
- If we are satisfied with its performance, we can now use this model to predict new data

Python

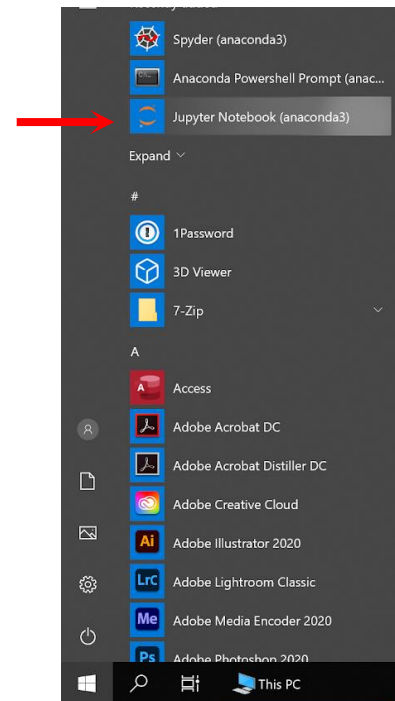
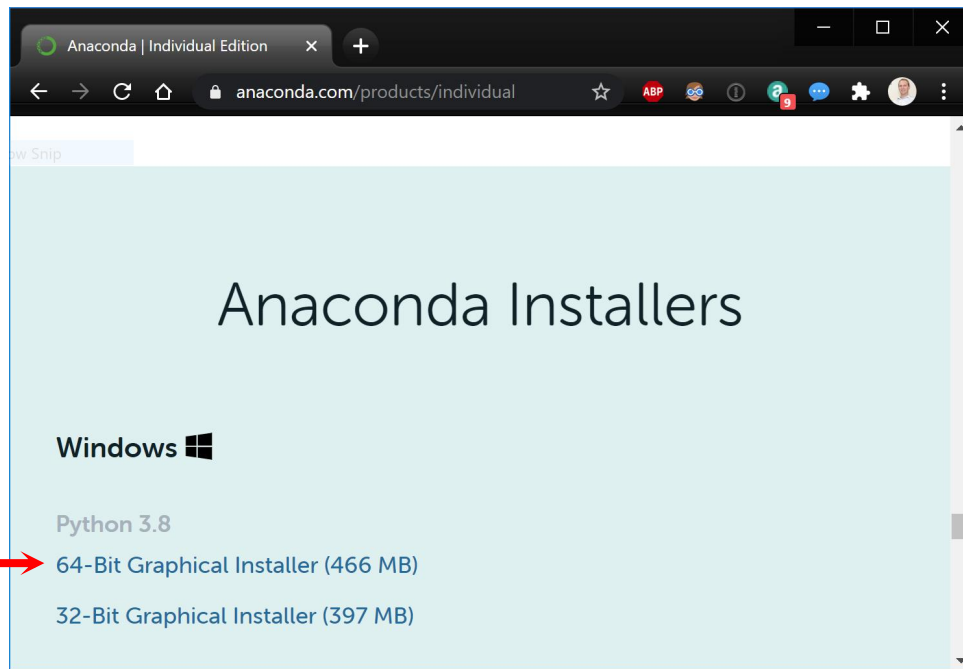
- We assume that you are familiar with the basics of python
 - Recommended textbook



Free PDF [here](#)


Programming Environment: Local

- In this course we are going to use
 - Python 3, NumPy, MatPlotLib, SciPy and Jupyter



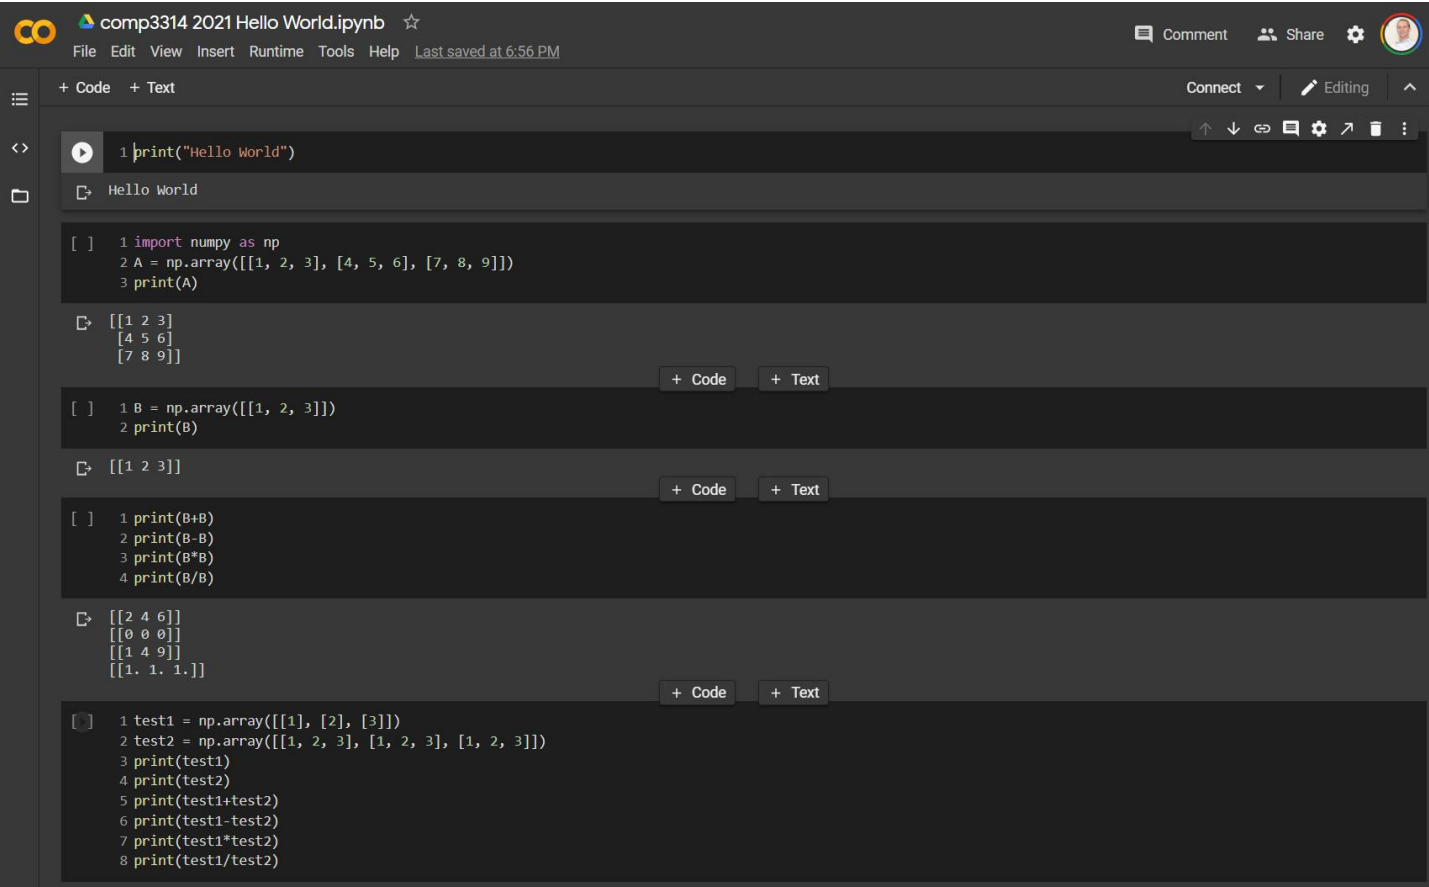
Programming Environment: Cloud

[Google CoLab](#)



Installation

→ Code is [here](#) ←



The screenshot shows a Jupyter Notebook titled "comp3314 2021 Hello World.ipynb". The interface includes a top bar with "File", "Edit", "View", "Insert", "Runtime", "Tools", and "Help" menus, and a status bar indicating "Last saved at 6:56 PM". The notebook contains four code cells, each with a "Run" button (play icon) and a "Clear" button (trash icon). The first cell contains a simple "Hello World" print statement. The second cell demonstrates creating a 2D NumPy array 'A' and printing it. The third cell creates a 1D array 'B' and prints it. The fourth cell performs various arithmetic operations on 'B' and prints the results. Each code cell is followed by its output, which is displayed in a separate box.

```
1 print("Hello World")
```

Hello World

```
[ ] 1 import numpy as np
    2 A = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])
    3 print(A)
```

```
[[1 2 3]
 [4 5 6]
 [7 8 9]]
```

```
[ ] 1 B = np.array([[1, 2, 3]])
    2 print(B)
```

```
[[1 2 3]]
```

```
[ ] 1 print(B+B)
    2 print(B-B)
    3 print(B*B)
    4 print(B/B)
```

```
[[2 4 6]
 [0 0 0]
 [1 4 9]]
[[1. 1. 1.]]
```

```
[ ] 1 test1 = np.array([[1], [2], [3]])
    2 test2 = np.array([[1, 2, 3], [1, 2, 3], [1, 2, 3]])
    3 print(test1)
    4 print(test2)
    5 print(test1+test2)
    6 print(test1-test2)
    7 print(test1*test2)
    8 print(test1/test2)
```

Linear Algebra Review


- We will only use basic concepts from linear algebra
- However, if you need a quick refresher, please take a look at Zico Kolter's excellent [videos](#)

Python Review

- We assume that you are familiar with the libraries/tools, follow these links if you need a refresher
 - [NumPy](#)
 - [Pandas](#)
 - [Matplotlib](#)
 - [Jupyter](#)

References

- Materials in this chapter are based on
 - [Book](#)
 - [Code](#)



References

- Some materials in this chapter are based on
 - [Book](#)
 - [Code](#)

